

Advanced Query Store and AQP in SQL Server 2016/2017



SQL Server 2016 / 2017 Mission-Critical (DB Engine)

Performance

Security

Availability / Platform

Scalability

Operational Analytics

Insights on operational data; Works with in-memory OLTP and disk-based OLTP

In-memory OLTP Enhancements

Greater T-SQL surface area, terabytes of memory supported, and greater number of parallel CPUs

Live Query Statistics

Query Store

Monitor and optimize query plans

Automatic Database Tuning



Provides insight into potential query performance problems, recommends solutions, and can automatically fix identified problems

DMV Improvements

Adaptive Query Processing



A feature family that introduces a new generation of query processing improvements

Always Encrypted

Sensitive data remains encrypted at all times with ability to query

Row-Level Security

Apply fine-grained access control to table rows

Dynamic Data Masking

Real-time obfuscation of data to prevent unauthorized access

Advanced Threat Detection

Ability to find unusual login patterns, track usage behavior in an auditing database, track SQL injection vulnerability, and more

Other Enhancements

Audit success/failure of database operations

TDE support for storage of in-memory OLTP tables

Enhanced auditing for OLTP with ability to track history of record changes

SQL Server 2017 on Linux



Enhanced AlwaysOn

Three synchronous replicas for auto failover across domains

Round robin load balancing of replicas

Automatic failover based on database health

DTC for transactional integrity across database instances with AlwaysOn

Support for SSIS with AlwaysOn

Stretch Database

Archive historical data transparently and securely to Azure

Queries stretch across local data as well as Azure data

Machine Learning Services



R Scripting along with Python scripting from the SQL Server Engine

Graph DB Support



For modeling many-to-many relationships

Enhanced Database Caching

Cache data with automatic, multiple TempDB files per instance in multi-core environments

New Programmatic Improvements

New TSQL Functionality, Maintenance Plan Improvements, New ALTER DATABASE Options

Expanded support for JSON data

New PolyBase query engine integrates SQL Server with external data in Hadoop or Azure Blob storage

Temporal Database Support

Query data as points in time


New in SQL Server 2017



SQL Server 2016/2017 Monitoring and Tooling

Full SQLTrace
Parity+ since
2012

Traditional Troubleshooting



Extended Events is scalable

Query Store is persisted and improving

Performance Dashboard Reports

Live Query Statistics

Lightweight Query Profiling

Expanded Query Plan Diagnostics

SSMS Dump Analysis (Preview)

SQL Server Vulnerability Assessment

SQL Data Discovery and Classification

"A Bad Plan is not the one which failed, but the one which succeeded at the Greatest Cost."

Anonymous DBA

Building Blocks: Plan Guides

How we got here..





Why Use Plan Guides?

- Useful for tuning queries generated by 3rd party applications
- Plan guides work by keeping a list of queries on the server, along with the Hints you want to apply
- You need to provide SQL Server with the query you want to optimize and a query hint using the OPTION clause
- When the query is optimized, SQL Server will apply the hint requested in the plan guide definition

Plan Guides Stored Procedures

- Use the `sp_create_plan_guide` stored procedure to create a plan guide
- Use `sp_control_plan_guide` to drop enable or disable plan guides
- You can see which plan guides are defined in your database using the `sys.plan_guides` catalog view
- **Note:** When Using Plan Guides, you must match Query Text and Parameter Names exactly

Common Query Hints Used in Plan Guides

- OPTIMIZE FOR (Value, Unknown)
- RECOMPILE
- MAXDOP #
- FORCE ORDER
- USE PLAN
- NULL

Plan Guides from Cache

- Also known as 'Plan Freezing'
- Prevents a current cached plan from changing
- Example:

```
-- Create a plan guide for the query by specifying the query plan in the plan
cache.
DECLARE @plan_handle varbinary(64);
DECLARE @offset int;
SELECT @plan_handle = plan_handle, @offset = qs.statement_start_offset
FROM sys.dm_exec_query_stats AS qs
CROSS APPLY sys.dm_exec_sql_text(sql_handle) AS st
CROSS APPLY sys.dm_exec_text_query_plan(qs.plan_handle,
qs.statement_start_offset,
qs.statement_end_offset) AS qp
WHERE text LIKE N'SELECT WorkOrderID, p.Name, OrderQty, DueDate%';

EXECUTE sp_create_plan_guide_from_handle
    @name = N'Guide1',
    @plan_handle = @plan_handle,
    @statement_start_offset = @offset;
GO

-- Verify that the plan guide is created.
SELECT * FROM sys.plan_guides
WHERE scope_batch LIKE N'SELECT WorkOrderID, p.Name, OrderQty, DueDate%';
GO
```

USE PLAN

- Used to explicitly guide the optimizer to use a specific plan
- Accepts an XML Showplan as the parameter
- **Note:** Plans larger than 8KB cannot be used

```
SELECT *
FROM Sales.SalesOrderHeader h, Sales.SalesOrderDetail
OPTION (USE PLAN N'<ShowPlanXML
xmlns="http://schemas.microsoft.com/sqlserver/2004/07/showplan"
Version="0.5"
Build="9.00.1187.07">
  <BatchSequence>
    <Batch>
      <Statements>
        ...
      </Statements>
    </Batch>
  </BatchSequence>
</ShowPlanXML>
')
```

.Now: Query Store

SQL Server 2016 / 2017
(Azure SQL DB)



Query and Query Plan Fingerprints

- Query Fingerprint
 - query_hash
 - Explicitly identifies a specific query in the cache.
 - *sys.dm_exec_requests*
 - *sys.dm_exec_query_stats*
- SQL Handle
 - sql_handle
 - Token for the SQL text that relates to a batch.
 - *sys.dm_exec_sql_text*
 - *sys.dm_exec_query_stats*
 - *sys.dm_exec_query_memory_grants*
- Query Plan Fingerprints
 - query_plan_hash
 - Useful to determine queries that share the same execution plan.
 - Can be used to determine if the query plan has changed.
 - *sys.dm_exec_requests*
 - *sys.dm_exec_query_stats*
- Plan Handle
 - plan_handle
 - Token for a cached execution plan.
 - *sys.dm_exec_query_plan*
 - *sys.dm_exec_cached_plans*

When performance is not good...

- Database is not working

Website / App is down




- Impossible to predict / root cause

Temporary Perf. issues



- Regression caused by upgrade

System Upgrade



Plan choice change can cause these problems

Have You Ever...?

...Had your system down/slowed down and everyone waiting for you to magically fix the problem ASAP?

...Upgraded and had an issue with the system down?

...Had a problem with performance and been unable to determine what was going on?

Query Plan
choice changes
can cause all of
these problems!

What are you doing today?

- Most solutions are reactive in nature
 - Flush the bad plan from the cache with `sp_recompile`
 - Flush the entire plan cache with `DBCC FREEPROCCACHE`
 - Force the plan to recompile every time
 - Restart OS / SQL Server (It works for some reason?)
- Proactive solutions are challenging
 - Often takes a long time to even detect there is a plan problem
 - Only the latest plan is stored in the cache
 - Need to catch both the good and the bad plan in order to troubleshoot
 - Information is stored in memory only
 - Reboot or memory pressure causes data to be lost
 - No history or timing available – stats are aggregated for what is currently in cache

Addressing Plan Choice Regressions

- First You Have to find the “Slow” Query
- Figuring out Why it is slow isn’t Easy
- You may not have enough information to fix it
- Even if you do know what it is supposed to be...
 - Can you modify the query to hint it?
 - Can you figure out how to make a plan guide?

Tackling the Problem – What Could We Do?

1. Store the history of plans for each query
2. Baseline the performance of each plan over time
3. Identify queries that have “gotten slower recently”
4. Find a way to force plans quickly and easily
5. Make sure this works across server restarts, upgrades, and query recompiles

This is what the Query Store does for you!

Introducing the Query Store

- Plan store persists execution plans per database
- Runtime stats store persists execution statistics per database
- New views and graphical interface allow you to quickly and easily troubleshoot query performance
 - Quickly find query plan performance regressions
 - Fix plan regressions by forcing a previous plan
 - Determine the number of times a query was executed in a given time window
 - Identify Top N Queries in the past X hours
 - Audit the history of query plans for a given query
 - Analyze the resource usage patterns for a particular database

Demonstration: Enabling Query Store in SQL Server 2016

- Query Store Properties



Key Usage Scenarios

Find and fix
query plan
regressions

Identify top
resource
consumers

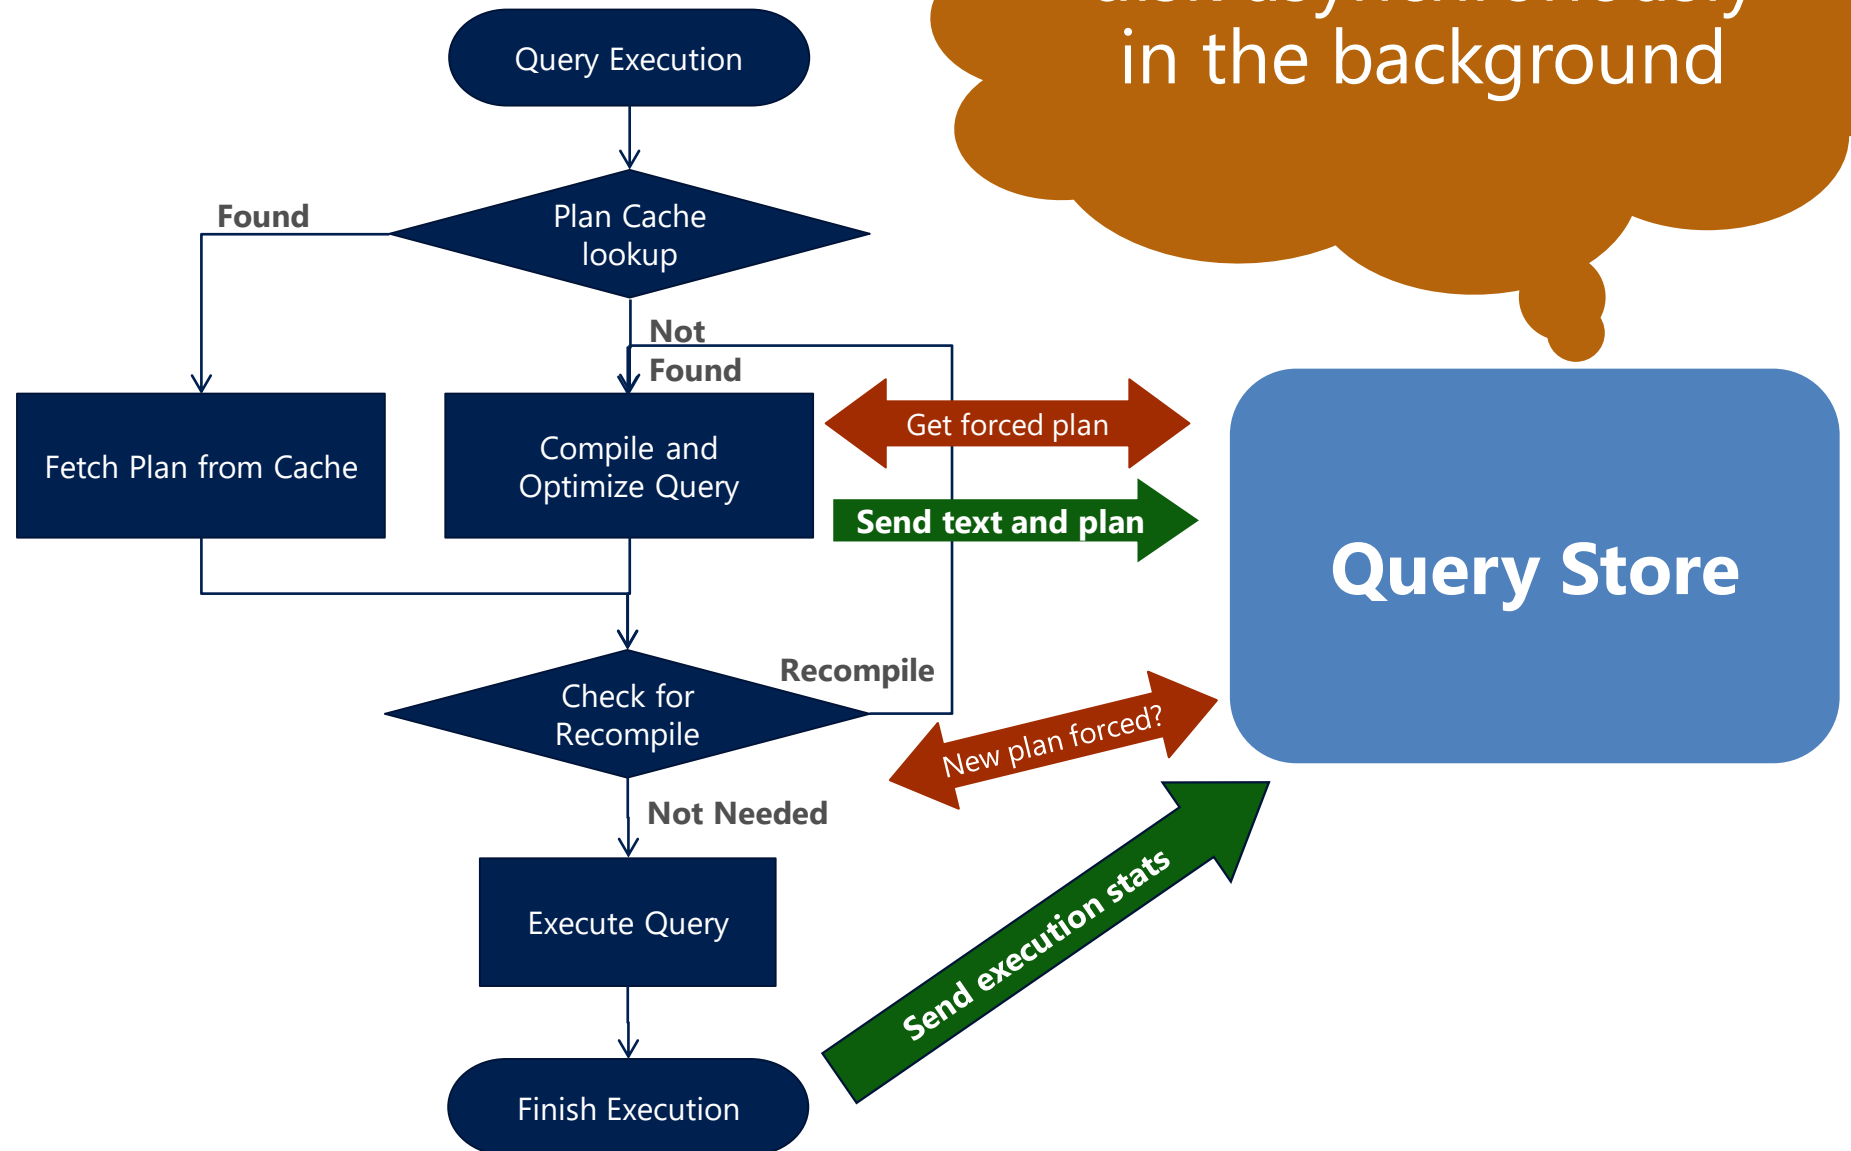
Reduce risks
with server
upgrade

Deep analysis
of workload
patterns/perf

Short-term/tactical

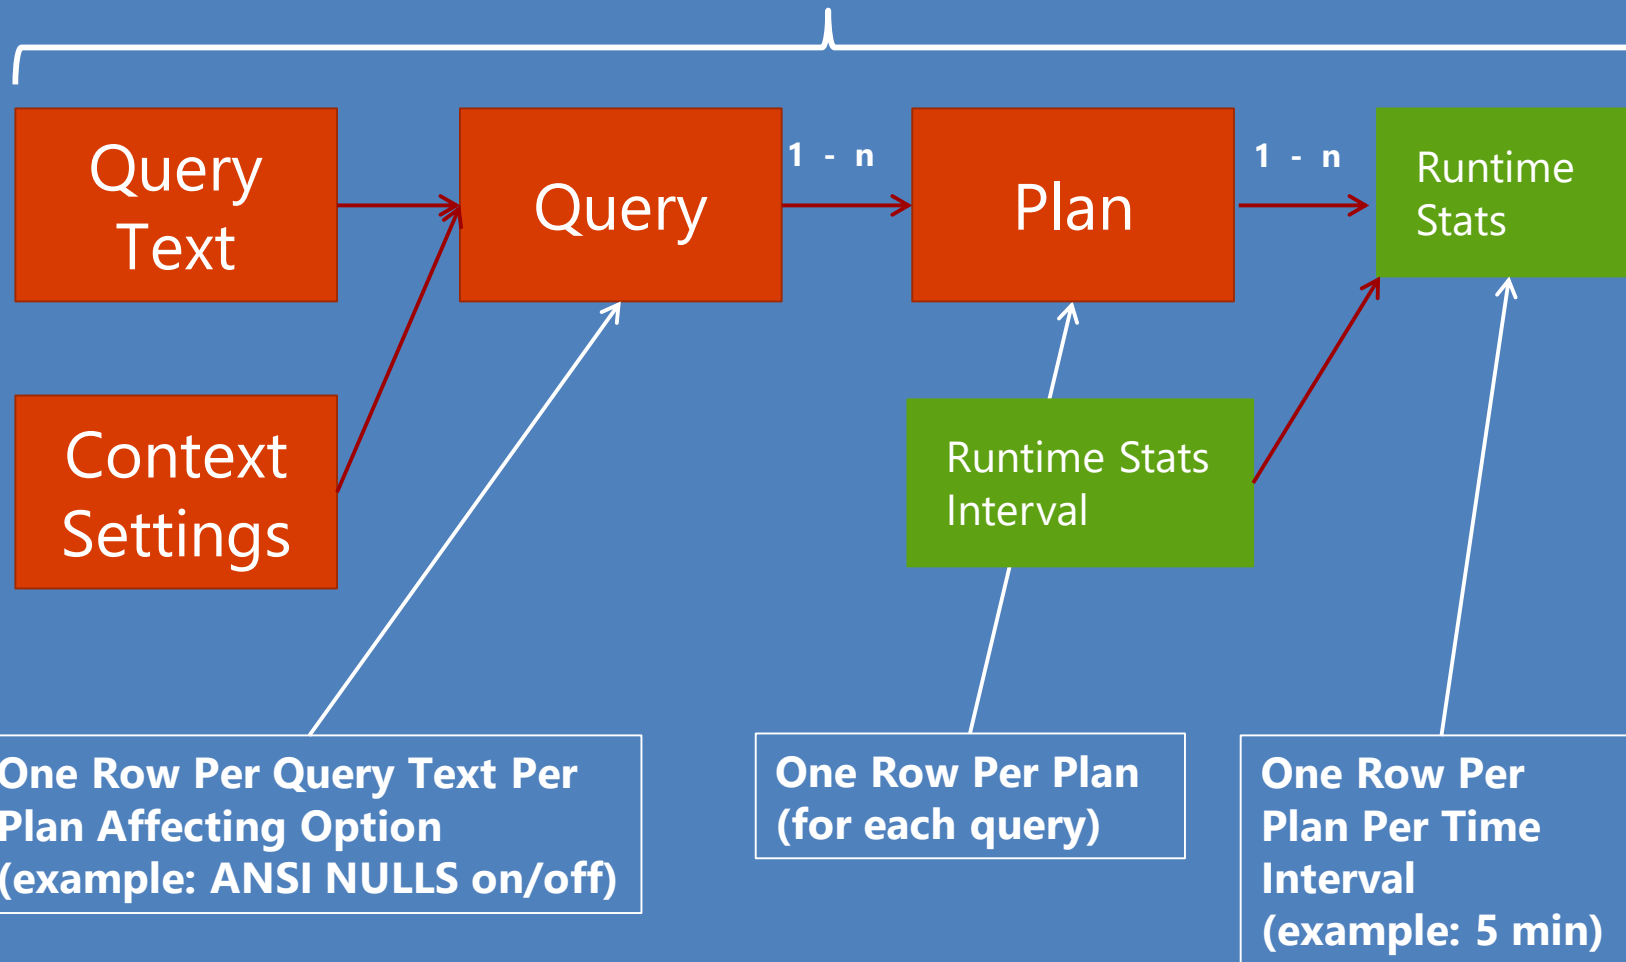
Long-term/strategic

SQL Query Execution

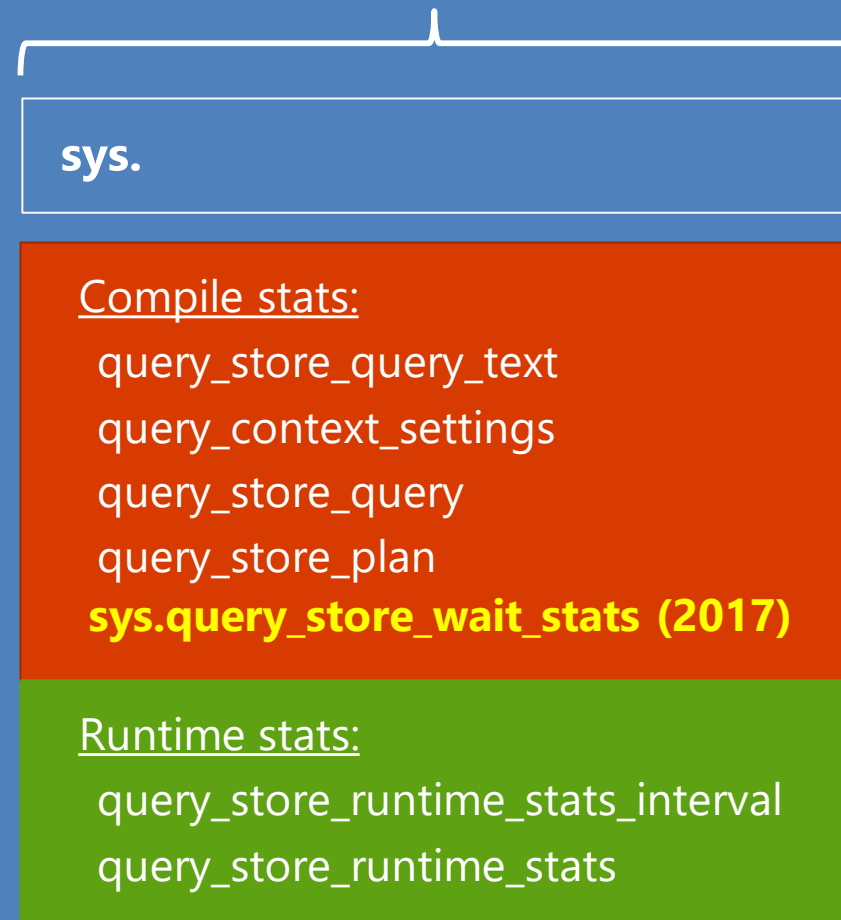


Query Store Schema Explained

internal tables



exposed views



Key DMVs for Query Store

```
SELECT * FROM sys.query_store_query_text
```

```
SELECT * FROM sys.query_store_query
```

```
SELECT * FROM sys.query_store_plan
```

```
SELECT * FROM sys.query_store_runtime_stats  
ORDER BY runtime_stats_id
```

```
SELECT * FROM  
sys.query_store_runtime_stats_interval
```

```
SELECT * FROM sys.query_store_wait_stats
```

```
SELECT * FROM sys.query_context_settings
```

- **The DMVs shown here are enabled and populated for each database when Query Store is turned on**

Query Store Details

- Plans and execution data are stored on disk in the user database
 - Query store data persists reboots, upgrades, restores etc.
 - Plans and statistics are tracked at the database level rather than the server level
- Query Store is configurable
 - Settings such as MAX_SIZE_MB, QUERY_CAPTURE_MODE, CLEANUP_POLICY allow you to decide how much data you want to store for how long
 - Can be configured either via the SSMS GUI or T-SQL scripts
- Query Store can be viewed and managed via scripting or SSMS

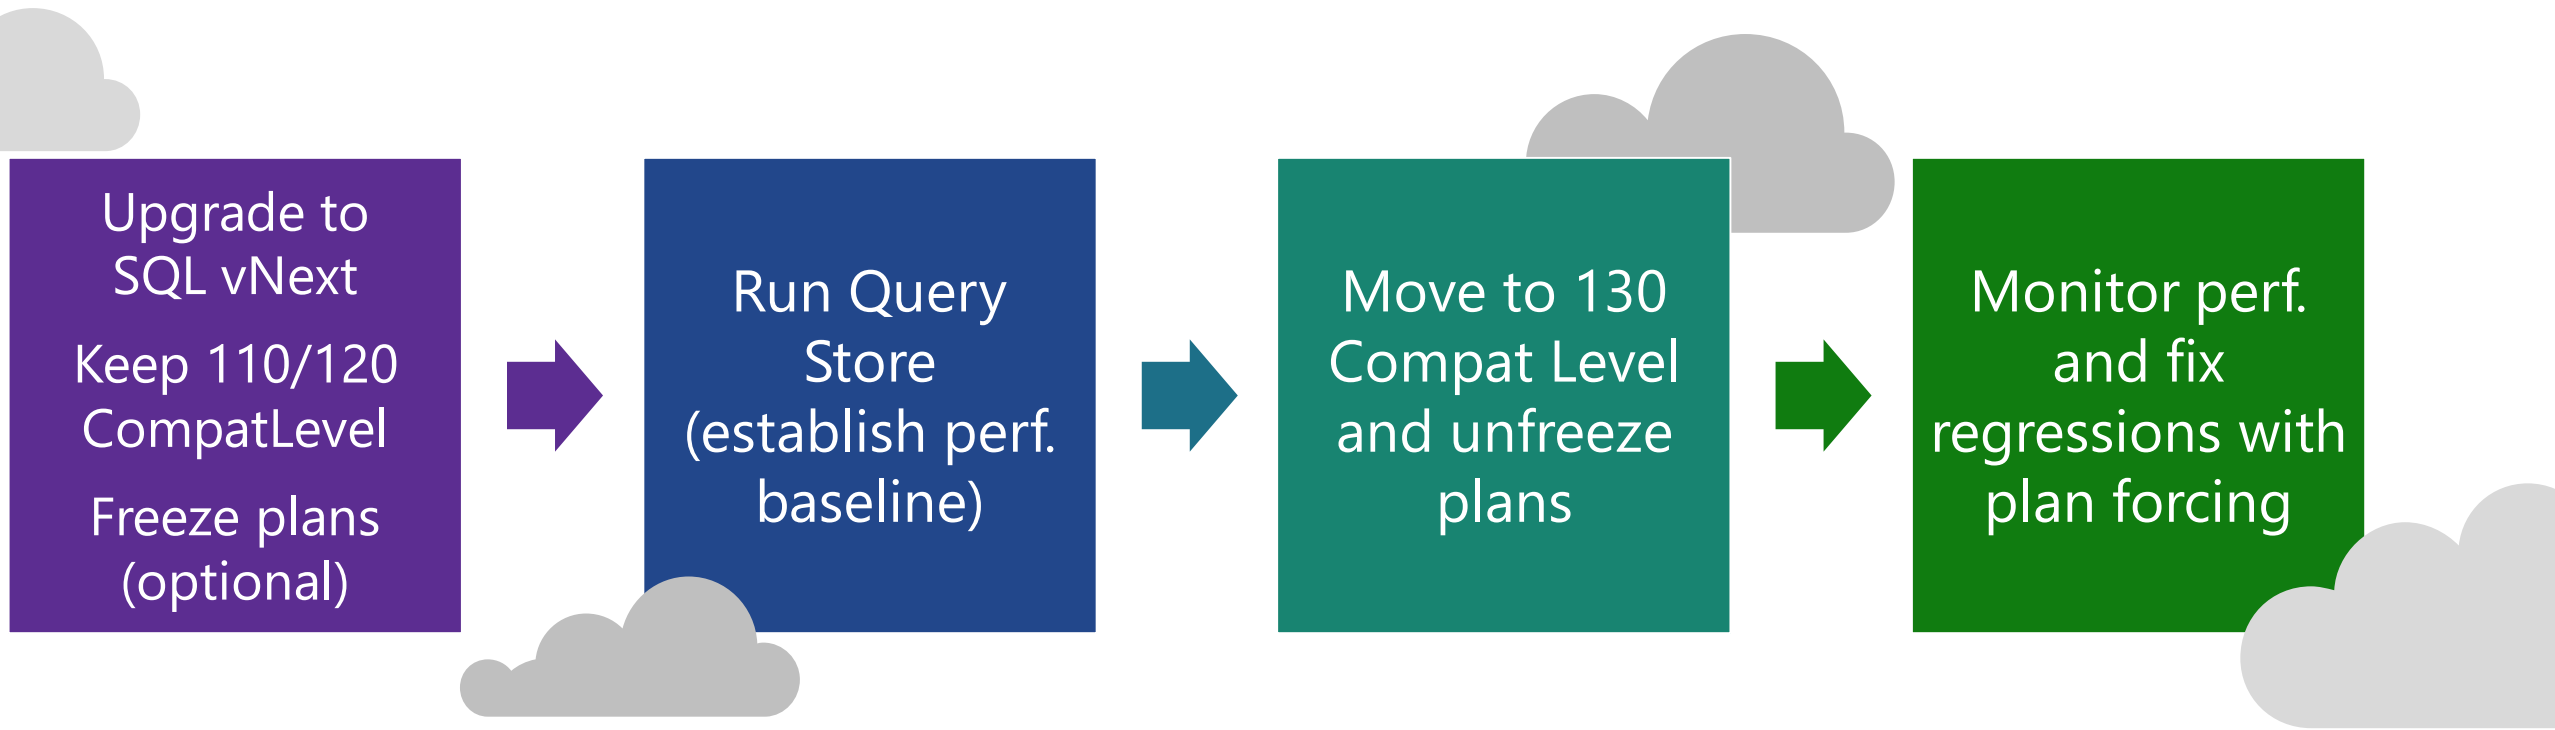
What does Query Store Track?

- Query Texts start at the first character of the first token of the statement; end at last character of last token
 - Comments before/after do not count
 - Spaces and comments inside *do* count
- Context_settings contains one row per unique combination of plan-affecting settings
 - Different SET options cause multiple "queries" in the Query Store
 - Plan caching/recompilation behavior unaffected

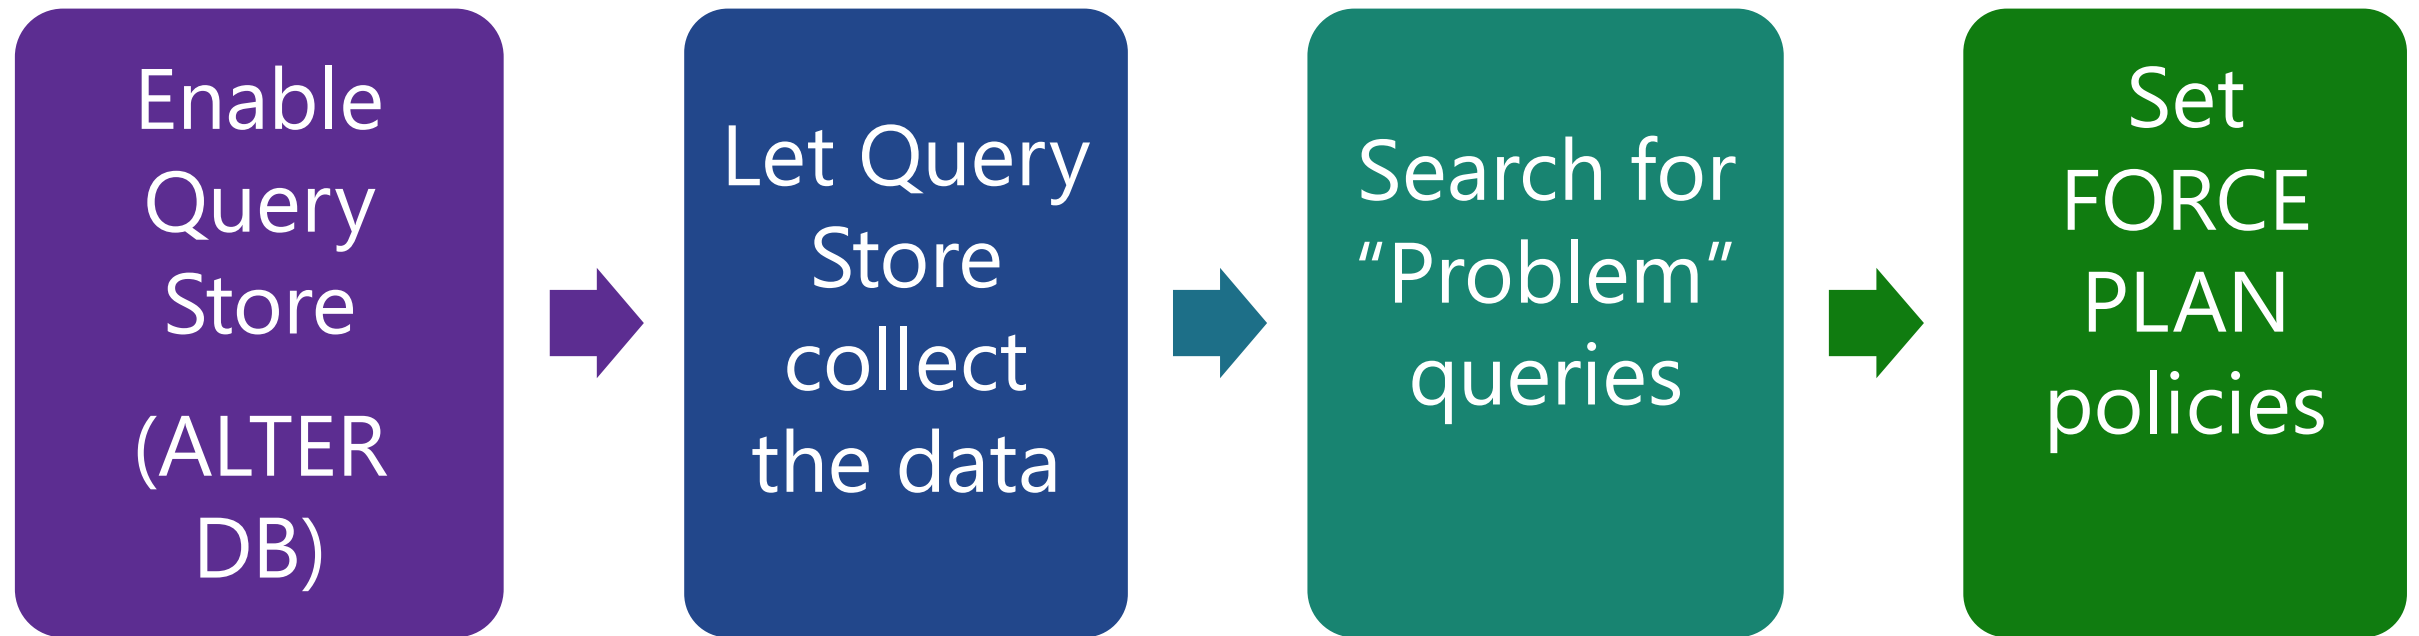
What Gets Captured?

- Query Texts
- Query Plans
- Runtime Statistics (per unit of time, default 1 hour)
 - Count of executions of each captured plan
 - For each metric: average, last, min, max, stddev
 - Metrics: duration, cpu_time, logical_io_reads, logical_io_writes, physical_io_reads, clr_time, DOP, query_max_used_memory, rowcount
 - Data is recorded when a query execution *ends*

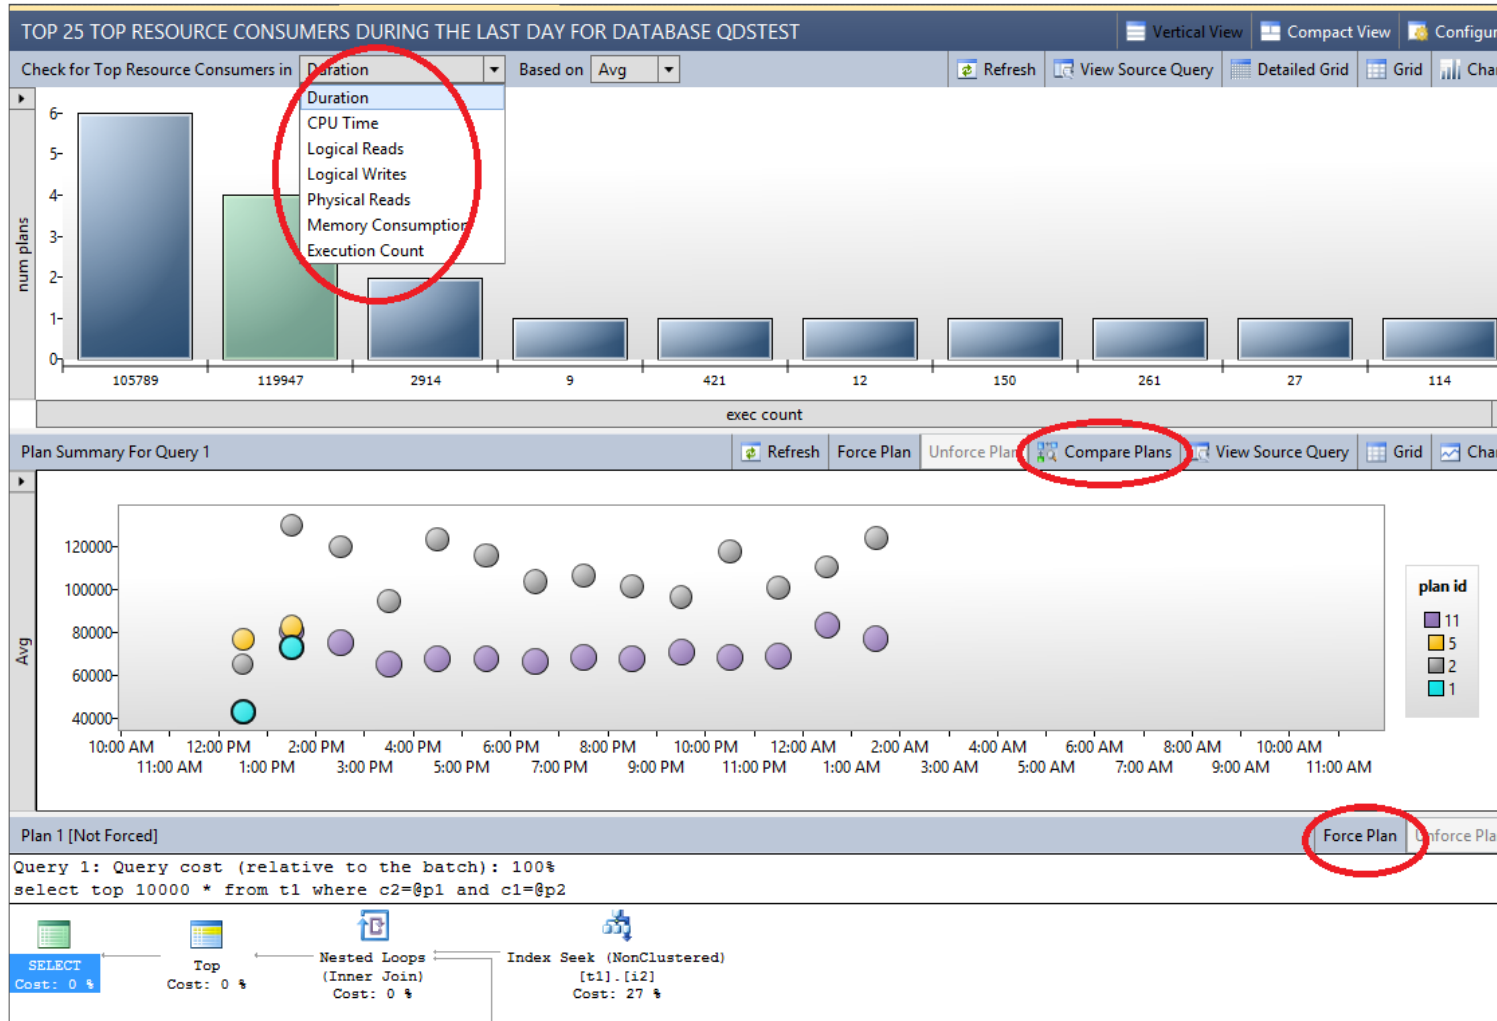
Keeping stability while upgrading to SQL Server 2016/2017



Troubleshooting with Query Store



Monitoring Performance with Query Store



- The Query Store feature provides DBAs with insight on query plan choice and performance

Working with Query Store

```
/* (6) Performance analysis using Query Store views*/
SELECT q.query_id, qt.query_text_id, qt.query_sql_text,
SUM(rs.count_executions) AS total_execution_count
FROM
sys.query_store_query_text qt JOIN
sys.query_store_query q ON qt.query_text_id =
q.query_text_id JOIN
sys.query_store_plan p ON q.query_id = p.query_id JOIN
sys.query_store_runtime_stats rs ON p.plan_id = rs.plan_id
GROUP BY q.query_id, qt.query_text_id, qt.query_sql_text
ORDER BY total_execution_count DESC

/* (7) Force plan for a given query */
exec sp_query_store_force_plan
12 /*@query_id*/, 14 /*@plan_id*/

);
```

```
/* (4) Clear all Query Store data */
ALTER DATABASE MyDB SET QUERY_STORE CLEAR;

/* (5) Turn OFF Query Store */
ALTER DATABASE MyDB SET QUERY_STORE = OFF;
```

- DB-level feature exposed through T-SQL extensions
- ALTER DATABASE
- Catalog views (settings, compile & runtime stats)
- Stored Procs (plan forcing, query/plan/stats cleanup)

Troubleshooting Query Store

- Plan forcing does not always work
 - **Example: If you drop an index, you can't force a plan that uses it.**
- Query Store will revert to not forcing if it fails
 - **This keeps the application working if the hint breaks**
- You can see which plans are failing to force by looking at the Plan Table:

```
SELECT * FROM sys.query_store_plan
WHERE is_forced_plan = 1 AND
force_failure_count > 0
```

Demonstration: Using Query Store in SQL Server 2016 / 2017



Queries with Forced Plans in SQL Server 2017

Query Store

Regressed Queries

Overall Resource Consumption

Top Resource Consuming Queries

Queries With Forced Plans

Queries With High Variation

Tracked Queries

Waits Reports Coming Soon!

Queries With Forced Plans

base AdventureWorks2016

plan id	plan count	last query execution time
2		2017-04-19 03:14:07.5030000
3		2017-04-19 02:56:51.9170000
2		2017-04-04 17:32:34.9870000
4		2017-04-26 20:39:48.3400000

Plan summary for query 469. Time period: Last month ending at 5/2/2017 11:11 PM

Plan Id

- 622
- 1795
- 1821

Force Plan

Unforce Plan

Query cost (relative to the batch): 100%

```
select p.object_id as object_id from ((select o.object_id as object_id, ct.depid as depid, sum(m.used_bytes) as used_memory_current from sys.objects$ o inner join sys.syssingleobjrefs c...
```

SELECT

Cost: 0 %

Filter

Cost: 0 %

Stream Aggregate (Aggregate)

Cost: 0 %

Nested Loops (Inner Join)

Cost: 1 %

Stream Aggregate (Aggregate)

Cost: 0 %

Nested Loops (Inner Join)

Cost: 0 %

Sort

Cost: 33 %

Nested Loops (Inner Join)

Cost: 0 %

Filter

Cost: 1 %

Table Valued Function [XTP_MM_CONSUMER]

Cost: 3 %

Index Seek (NonClustered) [sys.singleobjrefs].[nc1] [ct]

Cost: 44 %

Clustered Index Seek (Clustered) [sys.schobjs].[clst] [o]

Cost: 10 %

Filter

Cost: 0 %

Table Valued Function [XTP_MM_CONSUMER]

Cost: 5 %


Filter

Cost: 2 %

Demonstration: Query Store in SQL Server 2017

- Query Store Waits
- Queries with Forced Plans
- Queries with High Variations





SQL Server 2017 – Modern and Intelligent



Query Store – Wait Stats and “Cloud Learnings”

Automatic Tuning and Plan Correction

Query Plan Analysis in SSMS

Adaptive Query Processor

SQL Server 2017 – Query Store Improvements

- New Query Store Reports
- Automatic Tuning Feature Support

```
ALTER DATABASE AdventureWorks2017  
SET AUTOMATIC_TUNING ( FORCE_LAST_GOOD_PLAN = ON );
```

- DBCC CLONEDATABASE flushes statistics while cloning to avoid missing query store runtime statistics
- New DMVs
 - sys.query_store_wait_stats
 - sys.dm_db_tuning_recommendations
 - sys.database_automatic_tuning_mode
 - sys.database_automatic_tuning_options

SQL Server 2017 Automatic Tuning


```
ALTER DATABASE CURRENT
```

```
SET AUTOMATIC_TUNING (FORCE_LAST_GOOD_PLAN = ON);
```

Detect with **dm_db_tuning_recommendations** and force manually

Turn on Auto and system corrects

Reverts back to
"Last Known Good"



Perfect to help with [parameter sniffing](#)

Demonstration: Automatic Tuning in SQL Server 2017

- Query Store Waits
- Queries with Forced Plans
- Queries with High Variations



Adaptive Query Processing

A Family of Corrective Processing...



Risks of Misestimation



Slow Query
Response Time Due
to Bad Plans



Excessive Resource
Utilization
(CPU, Memory, IO)



Reduced Throughput
and Concurrency



T-SQL Refactoring
for Off-Model
Statements

Cardinality Estimation and Plan Quality

```
SELECT [fo].[Order Key], [fo].[Description], [fo].[Package], [fo].[Quantity],  
       [foo].[OutlierEventQuantity]  
FROM   [Fact].[OrderHistory] AS [fo]  
INNER JOIN [Fact].[WhatIfOutEventQuantity]('Mild Recession', '1-01-2013', '10-15-2014')  
        AS [foo] ON [fo].[Order Key] = [foo].[Order Key]  
                AND [fo].[City Key] = [foo].[City Key]  
                AND [fo].[Customer Key] = [foo].[Customer Key]  
                AND [fo].[Stock Item Key] = [foo].[Stock Item Key]  
                AND [fo].[Order Date Key] = [foo].[Order Date Key]  
                AND [fo].[Picked Date Key] = [foo].[Picked Date Key]  
                AND [fo].[Salesperson Key] = [foo].[Salesperson Key]  
                AND [fo].[Picker Key] = [foo].[Picker Key]  
INNER JOIN [Dimension].[Stock Item] AS [si]  
        ON [fo].[Stock Item Key] = [si].[Stock Item Key]  
WHERE [si].[Lead Time Days] > 0  
      AND [fo].[Quantity] > 50;
```

Adaptive Query Processing (SQL 2017)

Interleaved Execution

- Materialize estimates for multi-statement table valued functions (MSTVFs)
- Downstream operations will benefit from the corrected MSTVF cardinality estimate

Batch-mode Memory Grant Feedback

- Adjust memory grants based on execution feedback
- Remove spills and improve concurrency for repeating queries

Batch-mode Adaptive Joins

- Defer the choice of hash join or nested loop until after the first join input has been scanned
- Uses nested loop for small inputs, hash joins for large inputs

Demonstration: Adaptive Query Processing

- Interleaved Execution
- Batch-Mode Memory Grant Feedback
- Batch-Mode Adaptive Join



